

Dam Hong Phuc

+84 394240555 | damhongphuc.4104@gmail.com | [LinkedIn](#)

Education

FPT University, Ha Noi
B.Tech in Artificial Intelligence

Sep 2022 - Dec 2025
GPA: 8.49/10

Startup Product

Wizzdom

July 2025 – December 2025

Technical Leader & Founding Member

- Architected an AI-driven EdTech platform for Math educators that turns manual assessment into a 10x faster workflow with personalized student feedback.
- Spearheaded end-to-end development and rapid pivots toward product-market fit; engineered AI evaluation pipelines on **Gemini** and ran the cloud deployment on **AWS**.

Professional Experience

FPT Software

January 2026 – Present

AIOps & Platform Engineer

- Project: Enterprise HR Conversational AI Platform – Covestro AG**
 - Owned a live HR assistant (**LangGraph** multi-agent, **Bedrock Knowledge Base**, 1,000 articles, 9 languages) as sole developer, then **led its migration** from a monolith to an event-driven multi-service platform on **AWS** (Lambda, Bedrock, EventBridge, SQS, OpenSearch, S3, RDS, ElastiCache).
 - Established the first quality baseline with **RAGAS** + an independent LLM judge on Bedrock and a deterministic retrieval harness (252 questions); lifted recall@5 from 0.845 to 0.873 with a **config-only hybrid search + rerank** change and fixed an ingestion defect that flattened every table in the knowledge base.
 - Cut latency on repeat questions with a **semantic cache** (Cohere embeddings, ElastiCache, permission-scoped keys); migrated live chat from Azure GPT-4o to **Bedrock Claude Sonnet 4.5** with zero production errors via a PROD→DEV knowledge-base mirror later reused to restore 1,915 documents in 1h40.
 - Shipped an **HR-restricted knowledge base** with server-side, fail-closed permission filtering at a single choke point and a 4-stage rollout with no leakage window; instrumented **Langfuse** observability with PII dropped in production and real-time **PII redaction**.
 - Delivered a **zero-downtime production cutover** behind an SSM switch (rollback rehearsed at 34.8s; day one 528/528 turns, 0 errors) after proving quality parity 30/30 and cutting first-token latency 18.6s→6.3s; moved the stateful data plane between CDK apps via CloudFormation **stack refactor** + `cdk import` with zero id changes.
- Project: Enterprise AI Contract Intelligence System**
 - Upgraded the legacy system to a **Skill-based framework** to parse unstructured contracts. Employed prompt chaining and vector **RAG (Qdrant)** to slash hallucinations and boost extraction accuracy by 35%.
 - Routed multi-model traffic via **AWS Bedrock** with **SSE streaming**, and migrated the system from local Docker to enterprise **AWS** infrastructure to scale batch extraction to thousands of documents.

HTI Group

April 2025 – December 2025

AI Engineer & LLMops

- Project: Enterprise Agentic Chatbot & LLMops Infrastructure** 
 - Prototyped **Multi-Agent** solutions with **LangGraph** and **LlamaIndex (ReAct)** to unify institutional knowledge, deployed as cloud-native microservices on **AWS Lambda** and **ECS Fargate**.
 - Automated **CI/CD pipelines** via IaC (**AWS CDK/TypeScript**) for repeatable deployments. Implemented distributed observability with **Langfuse** to trace agent trajectories and ensure hallucination-free outputs.
 - Built a multi-format ingestion pipeline (Excel, Docs, PPT, PNG, CSV) with **Hierarchical Semantic Chunking** + LLM contextualization, and combined **Vector**, **Full-text**, **Graph** and **Pandas** retrieval with tool calling (Discord, Google Calendar).

HG Group

May 2023 – March 2025

AI Engineer

- Project: HG Chatbot - Enterprise Conversational AI**
 - Built parsing pipelines for PDF, Docx, PPT, Excel and PNG sources and a search engine with **Hybrid Retrieval** and **Cross-Encoder reranking** to maximize context relevance.
- Project: HG Render - AI Media Recommendation Engine**
 - Built a recommendation engine over 1M+ multimodal assets with dense retrieval (**Qdrant**) and **GPT**-driven combinatorial optimization to assemble production-ready videos.
- Project: HIG (HG Image Generation) - Internal AI Asset Creation**
 - Fine-tuned **FLUX.1** with LoRA to specific design styles and produced **custom visual assets**  (*Lofi-cartoon*) for enterprise media campaigns.

Technical Skills & Certifications

AI & Agentic Paradigms: Python, PyTorch, LangChain, LangGraph, LlamaIndex, FLUX.1 LoRA, vLLM

Distributed Systems: FastAPI, Event-driven Microservices, gRPC, Async Systems, Redis, Model Context Protocol

Databases & Messaging: PostgreSQL, MongoDB, OpenSearch, Qdrant, Kafka, RabbitMQ, AWS SQS, EventBridge

Cloud, DevOps & LLMops: AWS, Docker, Ray, K8s, Langfuse, LangSmith, GitHub Actions, CI/CD

Certifications: AWS SAA-C03, AWS AIP-C01, Project Manager, ML & DL Specialization, IBM Full Stack, IELTS 7.0